



Predict Heart Disease Using Logistic Regression

Import Libraries

```
In [1]: import joblib
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
```

```
In [2]: df = pd.read_csv("heart.csv")
df.head()
```

```
Out[2]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	0
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	0
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	0
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	0



Step 1: Explore the Dataset

```
In [3]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   age         1025 non-null   int64  
 1   sex         1025 non-null   int64  
 2   cp          1025 non-null   int64  
 3   trestbps    1025 non-null   int64  
 4   chol        1025 non-null   int64  
 5   fbs         1025 non-null   int64  
 6   restecg     1025 non-null   int64  
 7   thalach     1025 non-null   int64  
 8   exang       1025 non-null   int64  
 9   oldpeak     1025 non-null   float64 
10   slope       1025 non-null   int64  
11   ca          1025 non-null   int64  
12   thal        1025 non-null   int64  
13   target      1025 non-null   int64  
dtypes: float64(1), int64(13)
memory usage: 112.2 KB

```


In [4]: `df.describe()`

Out[4]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	target
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	54.434146	0.695610	0.942439	131.611707	246.000000	0.149268	0.529756	149.114146	0.336585	0.408642
std	9.072290	0.460373	1.029641	17.516718	51.59251	0.356527	0.527878	23.005724	0.472772	0.581717
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.000000
25%	48.000000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	132.000000	0.000000	0.000000
50%	56.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	152.000000	0.000000	0.000000
75%	61.000000	1.000000	2.000000	140.000000	275.000000	0.000000	1.000000	166.000000	1.000000	0.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	1.000000

🌸 Step 2: Exploratory Data Analysis

```
In [5]: plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()
```



Most Positively Correlated Features

```
In [6]: corr_matrix = df.corr()

target = 'target'

corr_with_target = corr_matrix[target].sort_values(ascending=False)

print("Most Positively Correlated Features:")
print(corr_with_target.head(6))
```

Most Positively Correlated Features:

target	1.000000
cp	0.434854
thalach	0.422895
slope	0.345512
restecg	0.134468
fbs	-0.041164

Name: target, dtype: float64

Most Negatively Correlated Features

```
In [7]: print("Most Positively Correlated Features:")
print(corr_with_target.tail(5))
```

Most Positively Correlated Features:

sex	-0.279501
thal	-0.337838
ca	-0.382085
exang	-0.438029
oldpeak	-0.438441

Name: target, dtype: float64

Step 3: Prepare the Data

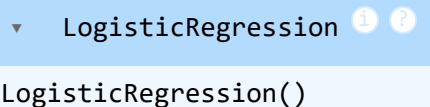
```
In [8]: # Split features and target
X = df[["cp", "thalach", "slope", "restecg", "fbs", "sex", "thal", "ca", "exang", "oldpeak"]]
y = df["target"]
```

```
In [9]: # Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [10]: # Feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Step 4: Train the Model

```
In [11]: # Train Logistic regression
model = LogisticRegression()
model.fit(X_train_scaled, y_train)
```

```
Out[11]: 
LogisticRegression()
```

Step 5: Evaluate the Model

```
In [12]: # Predict
y_pred = model.predict(X_test_scaled)
```

```
In [13]: # Evaluate
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)
```

```
In [14]: print("Accuracy:", accuracy)
print("\nClassification Report:\n", report)
```

Accuracy: 0.824390243902439

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.74	0.81	102
1	0.78	0.91	0.84	103
accuracy			0.82	205
macro avg	0.83	0.82	0.82	205
weighted avg	0.83	0.82	0.82	205



Step 6: Save the Model

```
In [15]: joblib.dump(model, 'heart_model.joblib')
```

```
Out[15]: ['heart_model.joblib']
```



Done! You've built a working heart disease predictor!